

# Chapter 7

Multicores,  
Multiprocessors, and  
Clusters



# Multithreading

- Performing multiple threads of execution in parallel
  - Replicate registers, PC, etc.
  - Fast switching between threads
- Fine-grain multithreading
  - Switch threads after each cycle
  - Interleave instruction execution
  - If one thread stalls, others are executed
- Coarse-grain multithreading
  - Only switch on long stall (e.g., L2-cache miss)
  - Simplifies hardware, but doesn't hide short stalls (eg, data hazards)



# Simultaneous Multithreading

- In multiple-issue dynamically scheduled processor
  - Schedule instructions from multiple threads
  - Instructions from independent threads execute when function units are available
  - Within threads, dependencies handled by scheduling and register renaming
- Example: Intel Pentium-4 HT
  - Two threads: duplicated registers, shared function units and caches



# Future of Multithreading

- Will it survive? In what form?
- Power considerations  $\Rightarrow$  simplified microarchitectures
  - Simpler forms of multithreading
- Tolerating cache-miss latency
  - Thread switch may be most effective
- Multiple simple cores might share resources more effectively



# Instruction and Data Streams

## ■ An alternate classification

		Data Streams	
		Single	Multiple
Instruction Streams	Single	<b>SISD:</b> Intel Pentium 4	<b>SIMD:</b> SSE instructions of x86
	Multiple	<b>MISD:</b> No examples today	<b>MIMD:</b> Intel Xeon e5345

## ■ SPMD: Single Program Multiple Data

- A parallel program on a MIMD computer
- Conditional code for different processors



# SIMD

- Operate elementwise on vectors of data
  - E.g., MMX and SSE instructions in x86
    - Multiple data elements in 128-bit wide registers
- All processors execute the same instruction at the same time
  - Each with different data address, etc.
- Simplifies synchronization
- Reduced instruction control hardware
- Works best for highly data-parallel applications



# Vector Processors

- Highly pipelined function units
- Stream data from/to vector registers to units
  - Data collected from memory into registers
  - Results stored from registers to memory
- Example: Vector extension to MIPS
  - $32 \times 64$ -element registers (64-bit elements)
  - Vector instructions
    - l v, sv: load/store vector
    - addv. d: add vectors of double
    - addvs. d: add scalar to each element of vector of double
- Significantly reduces instruction-fetch bandwidth



# Example: DAXPY ( $Y = a \times X + Y$ )

## ■ Conventional MIPS code

```
l.d    $f0, a($sp)      ; load scalar a
addiu  r4, $s0, #512     ; upper bound of what to load
loop:  l.d    $f2, 0($s0) ; load x(i)
      mul.d   $f2, $f2, $f0 ; a × x(i)
      l.d    $f4, 0($s1) ; load y(i)
      add.d   $f4, $f4, $f2 ; a × x(i) + y(i)
      s.d    $f4, 0($s1) ; store into y(i)
      addiu  $s0, $s0, #8 ; increment index to x
      addiu  $s1, $s1, #8 ; increment index to y
      subu   $t0, r4, $s0 ; compute bound
      bne    $t0, $zero, loop ; check if done
```

## ■ Vector MIPS code

```
l.d    $f0, a($sp)      ; load scalar a
lv     $v1, 0($s0)       ; load vector x
mulvs.d $v2, $v1, $f0    ; vector-scalar multiply
lv     $v3, 0($s1)       ; load vector y
addv.d $v4, $v2, $v3     ; add y to product
sv     $v4, 0($s1)       ; store the result
```





# Vector vs. Scalar

- Vector architectures and compilers
  - Simplify data-parallel programming
  - Explicit statement of absence of loop-carried dependences
    - Reduced checking in hardware
  - Regular access patterns benefit from interleaved and burst memory
  - Avoid control hazards by avoiding loops
- More general than ad-hoc media extensions (such as MMX, SSE)
  - Better match with compiler technology

